

ATIQUE AHMAD

Machine Learning & Full-Stack Software Engineer

engr.atique.ahmad@gmail.com | +92 324 0539099 | linkedin.com/in/atiqueahmad | github.com/atique-ahmad-01 | Lahore, Pakistan

Summary

Software engineer with 3+ years of experience across machine learning, LLMs and full-stack development. I train and deploy production ML models, build RAG and agentic AI systems, and ship end-to-end products with Python, Go, React and cloud infrastructure on AWS, GCP and Azure, backed by an MS in Data Science.

Experience

Software Engineer, ML/AI | *Paradigm Networks, Inc.*

Jan 2023 – Present

- **Factuality model:** Owned the full pipeline for a RAG factuality detector: curated 10K in-house samples, fine-tuned IBM Granite Guardian 3.1 (2B) with rsLoRA ($r=256$), a custom prompt-loss-weighted SFT trainer and class-balanced sampling on RunPod A40 GPUs, then merged the adapter and published weights to S3. Delivered **91.7% accuracy and 86.7% F1**, beating GPT-5 (86.3% accuracy) and Claude 3 Sonnet (74.9%).
- **Hallucination model:** Co-developed a two-phase fine-tuned Granite detector (64K open-source + 10K in-house samples) that reached **98.1% F1 and 97.5% accuracy**, ahead of GPT-5 (88.2% accuracy) and Gemini 2.0 Flash (85.0%); owned its production serving and upgrades.
- **Model serving & deployment:** Deployed both detectors as GPU microservices: vLLM's OpenAI-compatible server behind async FastAPI (logprob-based confidence, warm-up, connection pooling), packaged as multi-stage distroless Docker images with weights from S3, and shipped via CodeBuild/ECR and GitHub Actions to AWS EKS.
- **Agentic evaluation platform:** Architected and built the service (FastAPI, MongoDB, DeepEval G-Eval) that scores every AI-agent run from its trace across 9 metrics (task success, plan quality, workflow completion, hallucination, toxicity); added local judges matching GPT-5 on 89% of plan-quality verdicts, integrated across the Go control server, LangGraph platform and React web app.
- **RAG verification APIs:** Built chunk-level RAG response verification, answer relevancy (DeepEval **0.90 F1** vs. RAGAS 0.85) and DeepEval faithfulness checks, with multi-cloud LLM support across AWS Bedrock (Claude, Llama 2), GCP Vertex AI (Gemini), Azure OpenAI (managed identity) and OpenAI with provider fallback.
- **PR review agent:** Built a LangGraph code-review agent that analyses pull-request diffs and publishes findings, traced with LangSmith and benchmarked against CodeRabbit.
- **MLOps:** Built CI test pipelines (pytest, mocked MongoDB) and hardened ML containers against vulnerabilities.

Teaching Assistant / Lab Engineer | *ITU Lahore — mentored 100+ students*

Jan 2024 – Jul 2024

Software Engineer Intern | *Clicky — full-stack web development*

Jan 2022 – Oct 2022

Research & Projects

Retrieval-Augmented Reasoning for Small LLMs | *vLLM, FAISS, nomic/jina embeddings, Qwen3*

MSDS research

- Built a retrieval pipeline that turns Qwen3-8B teacher solutions (served with vLLM) into an embedding memory bank and a 600K-edge reasoning graph, then injects retrieved reasoning into a Qwen3-1.7B student.
- Gated retrieval raised held-out accuracy from **75.0% to 85.4% (+10.4 pp)**; step-level retrieval lifted MATH-500 from 70% to 76%. Reproduced Retrieval-of-Thought (arXiv:2509.21743) at +6.2 pp and proved through ablations that the injection site, not the templates, drives its failure mode.

Education

M.S. Data Science | *Information Technology University, Lahore*

2025 – Present

B.S. Computer Engineering | *Information Technology University, Lahore* / CGPA 3.23

2019 – 2023

Skills

Machine Learning:	Deep learning, LLM fine-tuning (LoRA/PEFT), NLP, computer vision, self-supervised learning
GenAI & Agents:	RAG, AI agents, prompt engineering, LLM evaluation, LangChain, LangGraph, LangSmith, DeepEval
ML & Data Tools:	PyTorch, Hugging Face, scikit-learn, XGBoost, FAISS, Pandas, NumPy, Tableau, Power BI
Backend:	Python, Go, FastAPI, Django REST Framework, REST APIs, MongoDB, PostgreSQL, SQL
Frontend:	React, TypeScript, JavaScript, Next.js, Tailwind CSS, HTML/CSS
MLOps & Cloud:	vLLM, Docker, Kubernetes (EKS), Helm, CI/CD (GitHub Actions), AWS, GCP, Azure, RunPod